

# Machine Learning Training Methods

*Field notes from an AI-coached study session*

Koushik Reddy Pasula · MS, Artificial Intelligence and Machine Learning  
AIML-500 · Indiana Wesleyan University

This artifact captures a guided study session in which I worked through the core machine-learning training methods in dialogue with an AI coach. Rather than reading definitions, I stated my own understanding of each concept first, tested it against the coach, and pushed back with follow-up questions. The value of the exercise was not memorizing terms — it was building an intuition I can explain to a non-specialist, and connecting each idea to a real engineering example.

## The learning approaches

### SUPERVISED LEARNING

#### How does a model learn from labeled data?

It learns from examples where the correct answer is attached. The model guesses, compares against the real answer, measures how far off it was, and adjusts to close the gap — repeated over many examples until it maps inputs to outputs.

*My anchor: studying with an answer key. Real example: fraud detection on card transactions, trained on past transactions already labeled fraud or not fraud.*

#### UNSUPERVISED LEARNING

##### What is the main approach when there are no labels?

No answer key exists. The model is handed raw data and finds structure on its own — grouping similar items, separating clusters, and flagging outliers, without being told what the groups mean.

*My anchor: customer segmentation — grouping users by behavior without predefining the segments. The honest test of quality is a downstream metric (does acting on the segments move a real number like click rate?), not just how clean the clusters look.*

#### REINFORCEMENT LEARNING

##### How does an agent learn the best actions?

Trial and error with rewards. The agent acts, receives a reward or penalty, and over many attempts learns which actions pay off — balancing exploring new moves against exploiting what already works. Delayed rewards are handled by *credit assignment*: when a win finally arrives, the reward is propagated backward across every step that led to it.

*My anchor: a restaurant tip. A big tip at the end teaches the server that the whole sequence — the greeting, the refills, the timing — contributed to the reward.*

## How training actually works

#### LOSS & OPTIMIZATION

##### How does "how wrong" become a direction to adjust?

A loss function scores how wrong each prediction is, and gradient descent uses that score to step the model's weights toward lower error. The optimization algorithm sets how big a step to take and which way — a poor one either crawls or overshoots the minimum.

*My anchor: walking downhill in fog. Loss is your altitude; the gradient tells you which way is down; you take a small step and measure again.*

#### ITERATION & OVERFITTING

##### How does repetition help, and when does it hurt?

One pass only nudges the model slightly; repeated passes let small corrections accumulate until the error settles. Too many passes and the model starts memorizing the training data instead of the pattern. The signal to stop is when validation error stops falling and starts rising.

*My anchor: overfitting is like hardcoding answers for the exact test cases you were given — it passes the exam and fails on real input. Holding back a test set is the only honest check that the model learned the pattern, not the answer key.*

#### ALGORITHMS & DATA

##### Why do algorithms and examples matter so much?

The algorithm is the method — it decides how patterns are interpreted and how the model updates. The data is the knowledge — it is the model's only teacher, so its quality, quantity, and variety set the ceiling on what can be learned. A wrong-fit algorithm still runs; it just produces a worse result.

*My anchor: using a hammer to drive a screw. The whole pipeline mirrors a software lifecycle — define the problem, gather and clean inputs, build, test on unseen cases, tune, deploy — where data preparation is the biggest hidden cost.*

## What the session came down to

Three ideas tie the entire topic together, and I can now state them in one line each:

- **The algorithm is the method** — how the model learns.
- **The data is the knowledge** — what the model can possibly learn.
- **Training is practice** — the loop of guessing, measuring error, and adjusting until the model generalizes to inputs it has never seen.